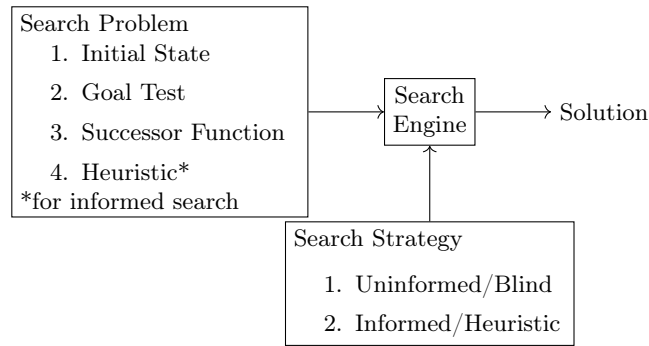


We can now expand on our earlier search problem model:



We will begin our discussion of heuristic searches by discussing best-first search. This is uninformed, but will lead us naturally into heuristics.


Uniform-Cost Search (UCS)

- uninformed generalization of BFS to weighted graphs
- expands based on contours of uniform cost and shape
- we goal check on expansion and not generation or our solution may be suboptimal
- we expand according to the function $g(n) := \text{cost of actions from initial state to node } n$

Properties:

1. complete? YES
2. optimal? YES
3. time? $O(b^{\lceil C^*/E \rceil})$
4. space? $O(b^{\lceil C^*/E \rceil})$

where E = minimal action cost & C^* = optimal solution cost



Unfortunately, this algorithm tends to wander a bit; we want a smarter algorithm!

Greedy Search

- a modified UFS based on an estimate of the distance to the goal state $h(n)$
- h is an **admissible** function $\equiv h(G) = 0$
- h is a **consistent** function $\equiv h(n) \leq \text{actual cost}$
This is stricter than the admissible requirement, but is hard to avoid in practice.

Properties:

1. complete? NO (for example: $A \text{ --- } 1 \text{ --- } B \text{ --- } 2 \text{ --- } C$)
2. optimal? NO
3. time? $O(b^m)$
4. space? $O(b^m)$

Let's compare our two algorithms:

UCS: $g(n)$: actual cost to get to n from initial state

- optimal
- conservative
- slow

Greedy: $h(n)$ estimate of cost to get to final state from n

- non-optimal
- aggressive
- fast

Neither of these has all the properties we want; can we get the best properties of both?

It turns out YES: we just have to add them directly!

$f(n) = g(n) + h(n)$ estimates the total cost & is admissible (provided $h(n)$ is admissible)

This leads us to a very important algorithm:

A* Search

- UCS based on $f(n) = g(n) + h(n)$
- forms contours which slim approaching goal
- prunes nodes outside contours
- we can prove that this algorithm is optimal:

Let's discuss the choice of heuristic: can we find a good heuristic for every problem?

$h(n) = 0$ works for every problem, but this makes the algorithm UCS!

Therefore, we want the maximum admissible heuristic for a given problem.

We will now consider a concrete example: 8 PUZZLE

We have two heuristics we would like to consider:

$h_1(n) = \#$ pieces out of place

$h_2(n) = \text{Manhattan Distance} \equiv \text{horizontal} + \text{vertical distance}$

$h_2(n) \leq h_1(n)$ for all n , so we say h_2 dominates h_1

Clearly h_2 is the better choice, but how much better can it really be?

We use two properties to evaluate search A^* search algorithms

1. effective branching factor (b^*) := avg branches per node
2. node count (N) = $1 + (b^*)^0 + (b^*)^1 + \dots + (b^*)^d$

d	IDS	$A^*(h_1)$	$A^*(h_2)$
4	$N = 112; b^* = 2.35$	$N = 13; b^* = 1.48$	$N = 12; b^* = 1.45$
6	$N = 680; b^* = 2.87$	$N = 20; b^* = 1.34$	$N = 18; b^* = 1.30$
8	$N = 6384; b^* = 2.73$	$N = 39; b^* = 1.33$	$N = 25; b^* = 1.24$
10	$N = 47,127$	$N = 93$	$N = 39$
12	$N = 3,644,035$	$N = 227$	$N = 73$

Clearly, A^* is much more efficient, and h_2 is much more efficient than h_1 ;

how do we formally evaluate this?

time cost = $O(b^\Delta)$, for absolute error $\Delta = h(n) - h^*$

\implies hypothetical min $O(h^\epsilon)$, for fractional error $\epsilon = (h - h^*)/h^*$ is the