

S	θ_S
male	0.55
female	0.45

S	C	$\theta_{c s}$
male	yes	0.05
male	no	0.95
female	yes	0.01
female	no	0.99

C	T_1	$\theta_{t_1 c}$
yes	+ve	0.8
yes	-ve	0.2
no	+ve	0.2
no	-ve	0.8

S	C	T_2	$\theta_{t_2 c,s}$
male	yes	+ve	0.80
male	yes	-ve	0.20
male	no	+ve	0.20
male	no	-ve	0.80
female	yes	+ve	0.95
female	yes	-ve	0.05
female	no	+ve	0.05
female	no	-ve	0.95

T_1	T_2	A	$\theta_{a t_1,t_2}$
+ve	+ve	yes	1
+ve	+ve	no	0
+ve	-ve	yes	0
+ve	-ve	no	1
-ve	+ve	yes	0
-ve	+ve	no	1
-ve	-ve	yes	1
-ve	-ve	no	0

Variables:

- C — condition
- T's — tests to detect the condition
- S — sex of the patient

We first compute the marginal distributions to discuss probabilities.
These come in two major forms:

T_1	0.2192
$\neg T_1$.7808

Prior Marginal
(pre-testing)

C	0.4533
$\neg C$.5467

Posterior Marginal
(assuming $T_1 = T_2 = 1$)

We then use this information to find the **Most Probable Explanation**:

We assume the final consequence (A) and search for the most probable query

Setting A moves us from $32 \rightarrow 8$ states;

result: $\{C=no, S=fe, T1=-ve, T2=-ve\} = 47\%$

This is not applicable in every situation; we generalize into the **Maximum a Posteriori Hypothesis**:

this is more complex and less efficient, but more often applicable

we find a subset of variables & find the MPE

(ex. $S=C \mid A \rightarrow \{C=no, S=ma\}$ | *approx* 49.3%)

When we seek to make inferences, algorithms fall into two major categories:

1. variable elimination
2. conditioning

In both situations, complexity is tied to the topology of the Bayesian Network

the actual property is **tree width** — it is roughly analogous to connectivity

$MPH = O(nd^w)$ given $n=\#var$, $d=\#val$, $w=width$

We will not define tree width, as it is complex, but we observe a few special cases:

Trees have width 1 (one path from any node to node)

Poly-trees (>1 parent ok) have width = the maximum parent count of any node

These are both singly connected networks; multiply connected leads to a DAG

Our first method of performing inference is

Weighted Model Counting

Consider the statement $\Delta = (A \vee B) \wedge \neg C$

This has 3 variables and suggests 8 worlds

Given an nd-DNNF circuit, we can solve in $O(N)$!

Consider the example on the right; it is trivial!


$WMC = 0.04 + 0.10 + 0.00 = 0.14$

We thus only need a compiler which would transform $\Delta \rightarrow$ sd-DNNF

Unfortunately, in the general case this is intractable

A	B	C	w
t	t	t	0.08
t	t	f	0.04
t	f	t	0.10
t	f	f	0.10
f	t	t	0.20
f	t	f	0.00
f	f	t	0.42
f	f	f	0.06

There are tractable subsets of this, though:



We can use probability as our weights! $\rightarrow WMC(\Delta \wedge \alpha) = Pr(\alpha)$

So we need to convert $\Delta \rightarrow$ boolean circuit.

We can see that Δ holds in $w_1, w_4, \&w_7$

We thus say $W(A) = W(\neg A) = \dots = W(\neg C) = 1$

We then let $W(P_i) = \theta_i \& W(\neg P_i) = 0$, so $W(A) = W(P_1)W(P_4)W(P_7)$

Thus we can solve probabilistic reasoning by symbol manipulations!

Properties of the algorithm:

1. there are many possible representations
2. it is not sensitive to tree width
3. it is not only applicable to Bayesian Networks

We can see that the $\text{size}(\text{Bayes}) = \text{size}(\text{CPT})$; though $\text{size}(\text{Bayes}) = O(nd^{k+1})$, $\text{size}(\text{joint-table}) = O(D^n)$!

This shows that Bayesian Networks are much more space efficient!

The process of modeling logic as a Bayesian Network has 3 steps:

1. define variables & values
2. define edges
3. specify CPT

Variables will then be labeled either query or evidence variables depending on the query.

This is a common approach used in early spam filters and Google ad Rephil.


Consider the following statements

- The cold causes a sore throat/chill
- The flu causes a sore throat/chill/fever/body ache
- Tonsillitis can show itself in fever/body ache

This network forms a bipartite graph.

We could have used one multivariable disease node.


We use this to give probabilities of a given condition.



If we have complete information, we can model a system as a Bayesian Network.

For incomplete information, however, we must use Expectation Maximization.

For example, suppose the following environment:



$P(P S) = 0.87$
$P(S \neg P) = 0.01$ & $P(\neg S P) = 0.1$
$P(U \neg P) = 0.1$ & $P(\neg U P) = 0.3$
$P(B \neg P) = 0.1$ & $P(\neg B P) = 0.2$

$P(p)$	0.87
$P(\neg p)$	0.13

S	P	$P(S, P)$
1	1	0.9
1	0	0.1
0	1	0.01
0	0	0.99

What is the marginal $(P|\neg S, \neg B, \neg U)$? – 10.21%!

WHAT?? why is that so high?

It turns out this is because of the false negative rate of the scanning test.

We want a false negative rate of below 5%.

We can address this in one of three ways:

1. get a better scanning test — a false (-) for S of 4.63% meets our requirement
2. lower the success of the procedure — 75.59% would meet our requirement
3. increase $P(L|P)$ — 99.67% meets our requirement

(b) and (c) turn out to be either impractical or in-economical, so our standard approach is to pay!